



**NTNU – Trondheim**  
Norwegian University of  
Science and Technology

# **MPI Programming**

Henrik R. Nagel  
Scientific Computing  
IT Division

# Outline

- Introduction
- Finite Difference Method
- Finite Element Method
- LU Factorization
- SOR Method
- Monte Carlo Method
- Molecular Dynamics
- MPMD Models



# Introduction

- OpenMP
  - can “easily” be added to existing serial programs
  - 1 node only
  - Semi-automatic parallelization
    - High communication bandwidth (memory)
    - Critical variables slow down performance
  - Limited amount of memory available (GB)
- MPI
  - should be use from the start
  - as many nodes as needed
  - Manual parallelization
    - Slow communication bandwidth (network)
    - Programmer has full control
  - All memory is available (TB)



# Used MPI Calls

- MPI\_SEND, MPI\_RECV
- MPI\_ISEND, MPI\_Irecv, MPI\_WAITALL
- MPI\_GATHER
- MPI\_BCAST, MPI\_REDUCE, MPI\_ALLREDUCE



# Getting Started

- Based on:
  - <http://www.redbooks.ibm.com/abstracts/sg245380.html>
- Kongull:
  - `ssh -Y kongull.hpc.ntnu.no`
  - `cp -r /home/hrn/Kurs/mpi .`
  - `module load intelcomp`
- Vilje:
  - `ssh -Y vilje.hpc.ntnu.no`
  - `cp -r /home/ntnu/hrn/Kurs/mpi .`
  - `module load intelcomp`



# Finite Difference Method

- Example 1 in the source code
- Only a skeleton of a 2D FDM is shown here
- Coefficients and the enclosing loop are omitted
- Data dependencies exist in both dimensions



# The Sequential Algorithm

```
PROGRAM main
IMPLICIT REAL*8 (a-h,o-z)
PARAMETER (m = 6, n = 9)
DIMENSION a(m,n), b(m,n)
DO j = 1, n
  DO i = 1, m
    a(i,j) = i + 10.0 * j
  ENDDO
ENDDO
DO j = 2, n - 1
  DO i = 2, m - 1
    b(i,j) = a(i-1,j) + a(i,j-1) + a(i,j+1) + a(i+1,j)
  ENDDO
ENDDO
END
```



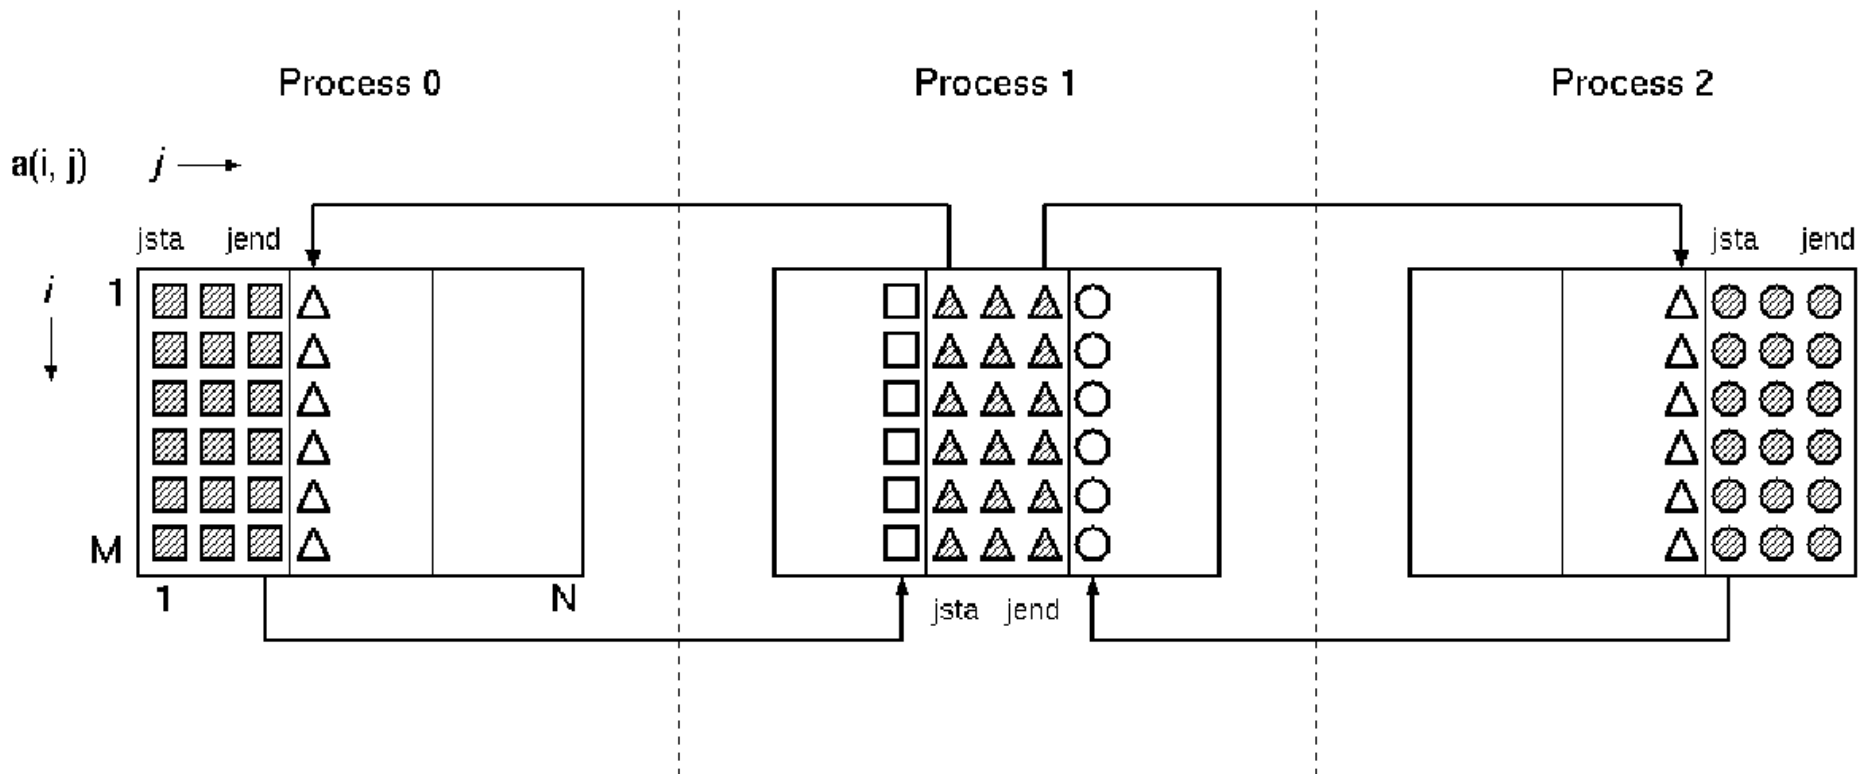
# Column-Wise Block Distribution

- Code example: ex1/fdm1.f90
- We must distribute a 2D matrix onto the processes
- Fortran stores arrays in column-major order
- Boundary elements between processes are contiguous in memory
- There are no problems with using MPI\_SEND and MPI\_RECV





# Column-Wise Block Distribution

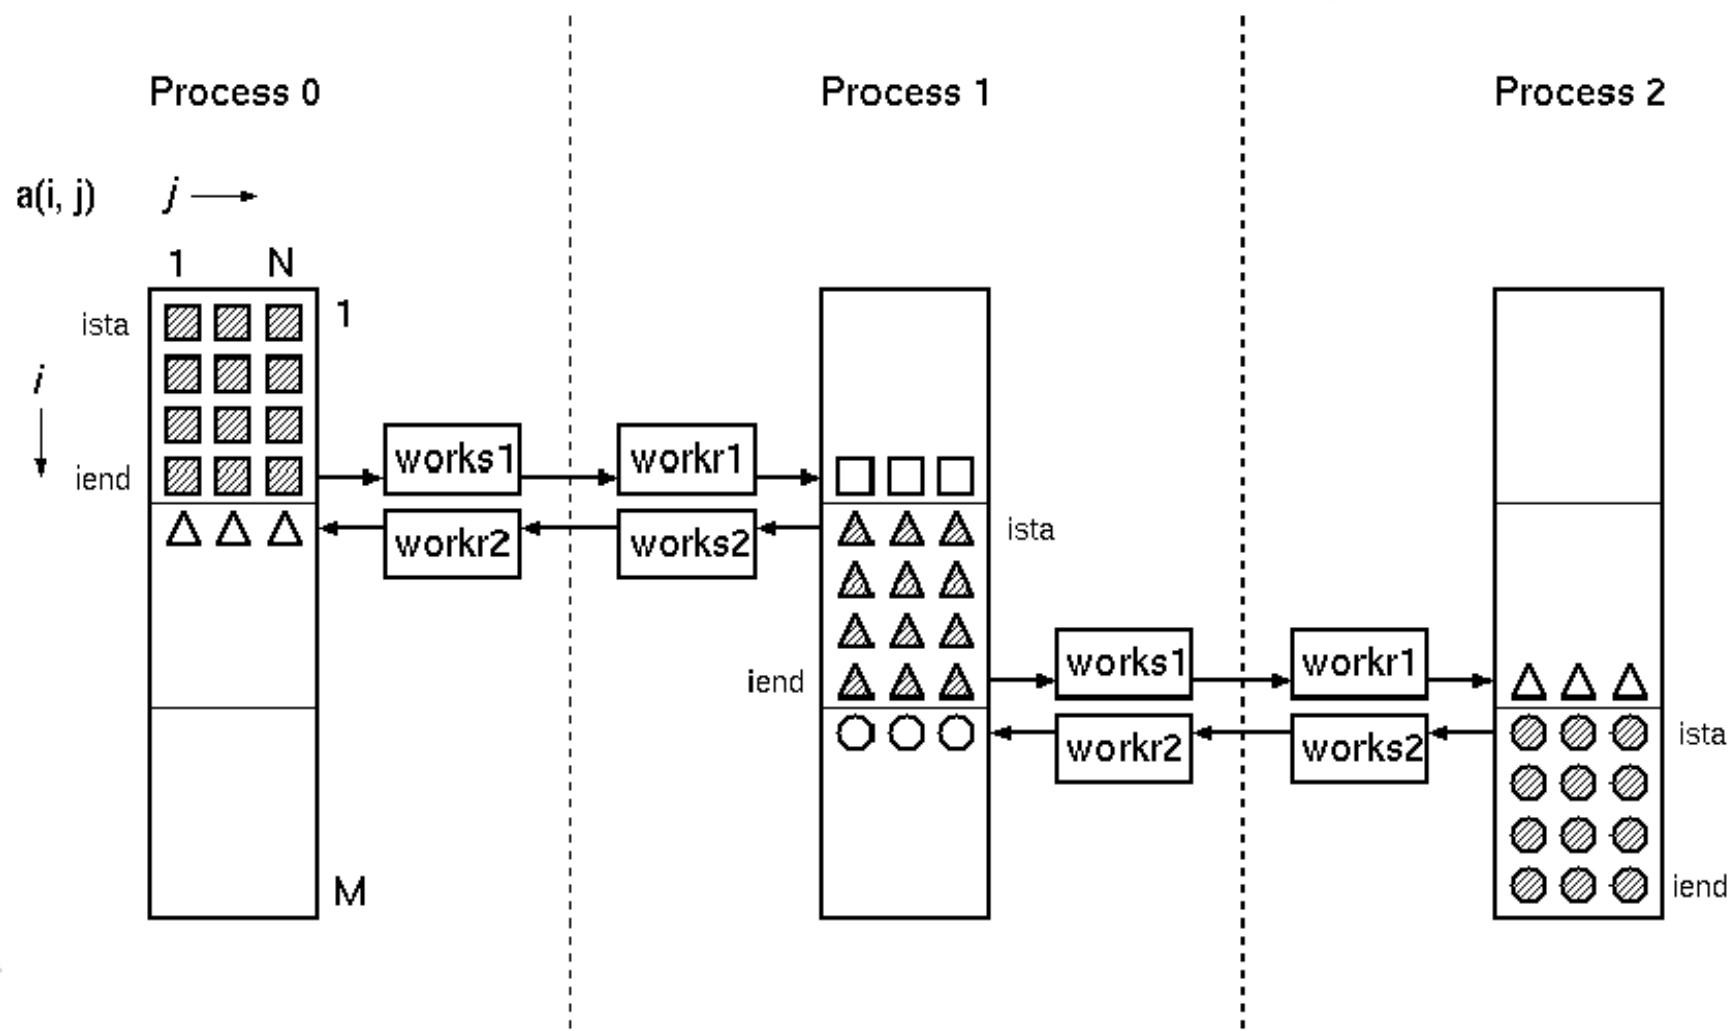


# Row-Wise Block Distribution

- Code example: `ex1/fdm2.f90`
- Fortran stores arrays in column-major order
- Boundary elements between processes are not contiguous in memory
- Boundary elements can be copied by:
  - Using derived data types
  - Writing code for packing data, sending/receiving it, and then unpacking it



# Row-Wise Block Distribution



# Block Distribution in Both Dim. (1)

- Code example: `ex1/fdm3.f90`
- The amount of data transferred might be minimized
  - Depends upon the matrix size and the number of processes
- A process grid itable is prepared for looking up processes quickly



# Block Distribution in Both Dim. (1)


$a(i, j)$      $j \rightarrow$

		1								N
$i \downarrow$	1	0	0	0	1	1	1	2	2	2
		0	0	0	1	1	1	2	2	2
		0	0	0	1	1	1	2	2	2
		0	0	0	1	1	1	2	2	2
		3	3	3	4	4	4	5	5	5
		3	3	3	4	4	4	5	5	5
		3	3	3	4	4	4	5	5	5
		3	3	3	4	4	4	5	5	5
		6	6	6	7	7	7	8	8	8
		6	6	6	7	7	7	8	8	8
		6	6	6	7	7	7	8	8	8
M		6	6	6	7	7	7	8	8	8

(a) The distribution of  $a()$

			$j \rightarrow$					
				-1	0	1	2	3
$i \downarrow$	-1	null	null	null	null	null	null	null
	0	null	0	1	2	null	null	null
	1	null	3	4	5	null	null	null
	2	null	6	7	8	null	null	null
	3	null	null	null	null	null	null	null

(b) The process grid



# Block Distribution in Both Dim. (2)

- Code example: `ex1/fdm4.f90`
- The corner elements are now included
- The data dependencies are therefore more complex




# The Sequential Algorithm

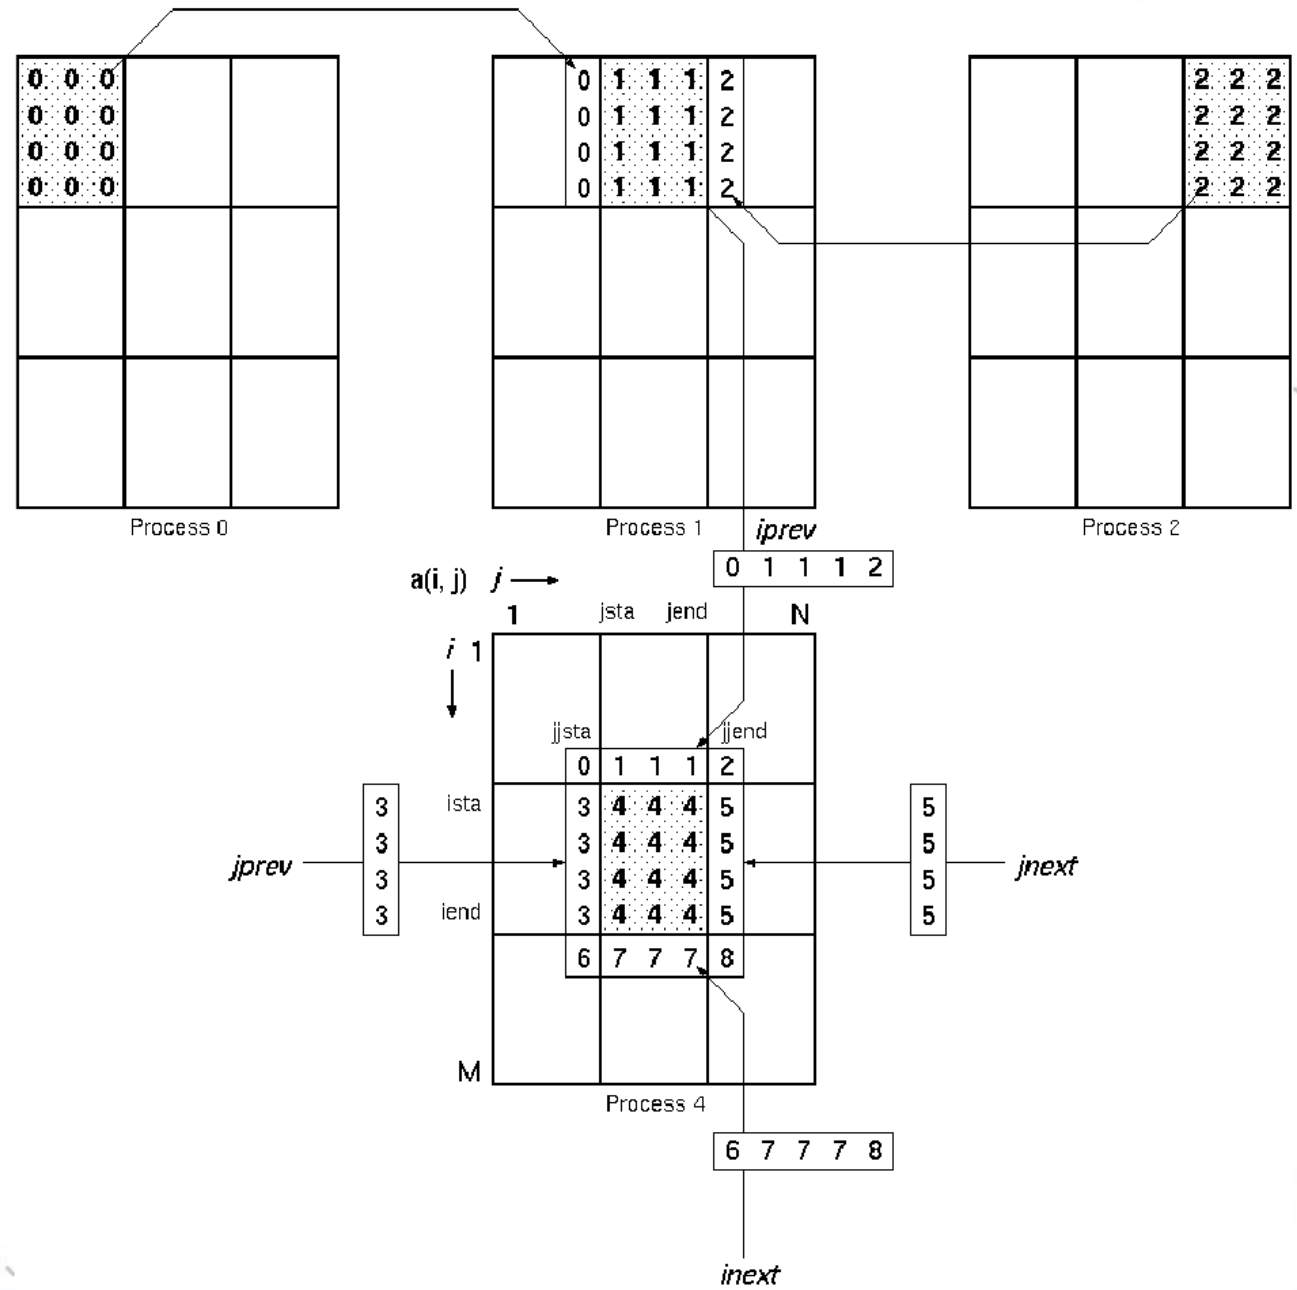
```
PROGRAM main
IMPLICIT REAL*8 (a-h,o-z)
PARAMETER (m = 12, n = 9)
DIMENSION a(m,n), b(m,n)
DO j = 1, n
  DO i = 1, m
    a(i,j) = i + 10.0 * j
  ENDDO
ENDDO
DO j = 2, n - 1
  DO i = 2, m - 1
    b(i,j) = a(i-1,j ) + a(i, j-1) + a(i, j+1) + a(i+1,j ) + &
      a(i-1,j-1) + a(i+1,j-1) + a(i-1,j+1) + a(i+1,j+1)
  ENDDO
ENDDO
END
```





# The Data Dependency






# Finite Element Method

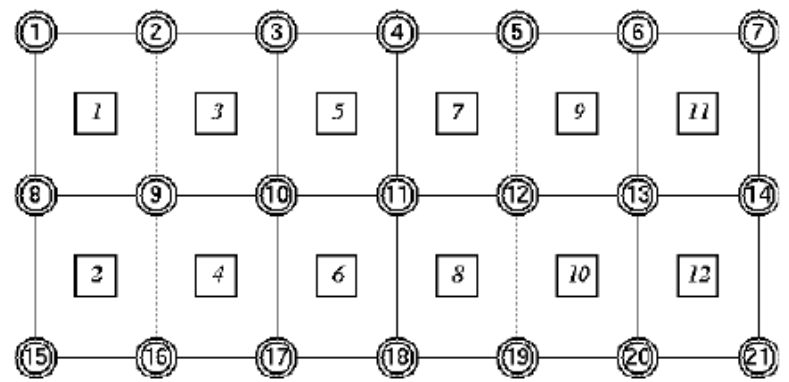
- Example 2 in the source code
- A more complete example that produces a result




# Finite Element Method




(a) Elements -> Nodes



(b) Update Nodes



(c) Nodes -> Elements



(d) Update Elements

m  
y of  
ogy

# The Sequential Algorithm

```

PARAMETER(iemax = 12, inmax = 21)
REAL*8 ve(iemax), vn(inmax)
INTEGER index(4,iemax)
...
DO ie = 1, iemax
    ve(ie) = ie * 10.0
ENDDO
DO in = 1, inmax
    vn(in) = in * 100.0
ENDDO
DO itime = 1, 10
    DO ie = 1, iemax
        DO j = 1, 4
            vn(index(j,ie)) = vn(index(j,ie)) + ve(ie)
        ENDDO
    ENDDO
ENDDO

```

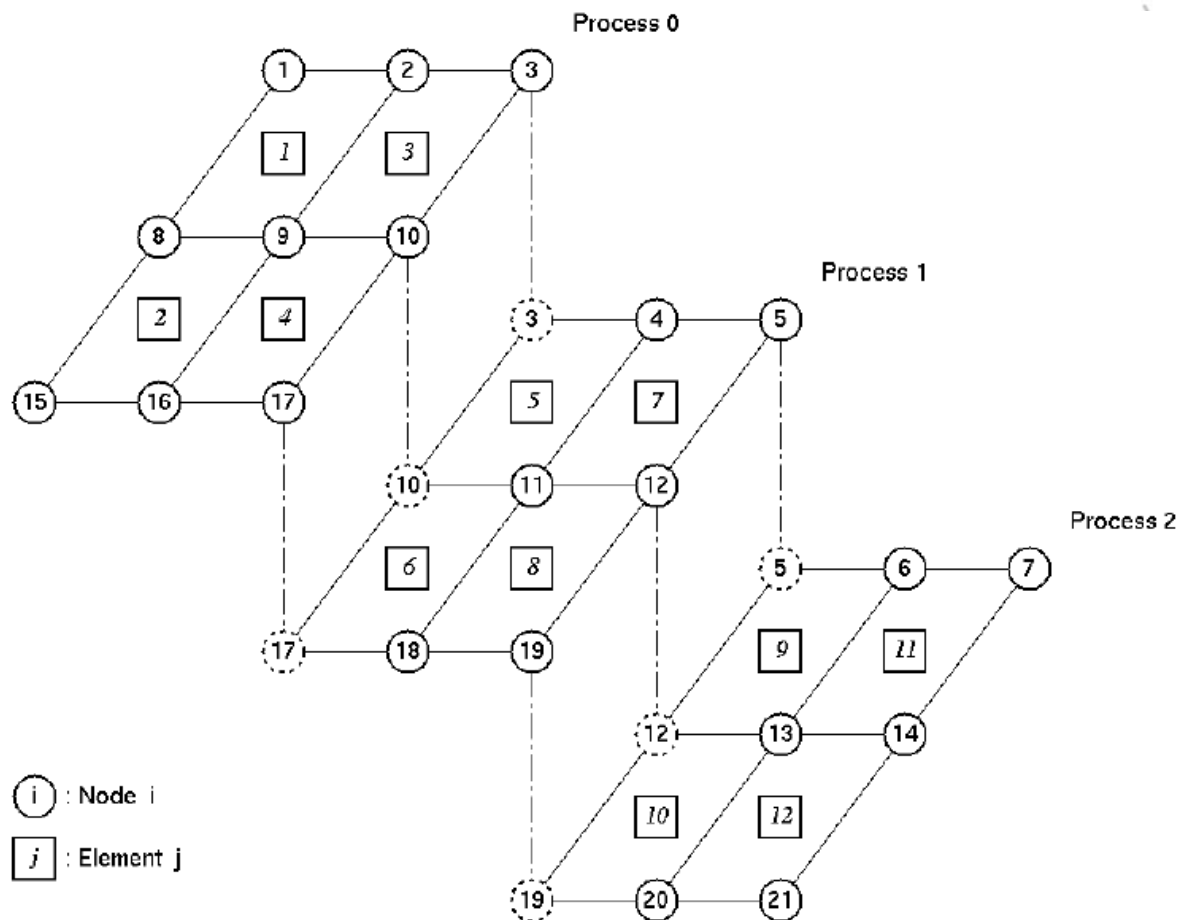
```

DO in = 1, inmax
    vn(in) = vn(in) * 0.25
ENDDO
DO ie = 1, iemax
    DO j = 1, 4
        ve(ie) = ve(ie) + vn(index(j,ie))
    ENDDO
ENDDO
DO ie = 1, iemax
    ve(ie) = ve(ie) * 0.25
ENDDO
ENDDO
PRINT *, 'Result', vn, ve

```



# Distributing the Data



# The Parallel Code

- Example ex2 in the source code
- Differences from original IBM version
  - 2D enumeration (row, column) is used instead of 1D enumeration
  - The amount of memory allocated by each process is minimized
  - A node column is sent to the right
  - An element column is sent to the left



# LU Factorization

- Example 3 in the source code
- Actually just Gaussian Elimination
- Solving a system of linear equations:  $Ax = b$
- Parallel ESSL has subroutines for LU factorization (outside the scope of this MPI course)
- Pivoting and loop-unrolling is not considered





# The Sequential Algorithm

```

PROGRAM main
PARAMETER (n = ...)
REAL a(n,n), b(n)
...
! LU factorization
DO k = 1, n-1
  DO i = k+1, n
     $a(i,k) = a(i,k) / a(k,k)$ 
  ENDDO
  DO j = k+1, n
    DO i = k+1, n
       $a(i,j) = a(i,j) - a(i,k) * a(k,j)$ 
    ENDDO
  ENDDO
ENDDO

```


```



! Forward elimination
DO i = 2, n
  DO j = 1, i - 1
     $b(i) = b(i) - a(i,j) * b(j)$ 
  ENDDO
ENDDO
! Backward substitution
DO i = n, 1, -1
  DO j = i + 1, n
     $b(i) = b(i) - a(i,j) * b(j)$ 
  ENDDO
   $b(i) = b(i) / a(i,i)$ 
ENDDO
...
END

```



# Distributing the Data






# The SOR Method

- Example 4 in the source code
- Solving the 2D Laplace equation using the Successive Over-Relaxation (SOR) method



# The Basic SOR Method



- : Updated in the current iteration
- ◐ : About to be updated
- : Not updated yet in the current iteration




# The Sequential Algorithm

```
PROGRAM sor
PARAMETER (mmax = 6, nmax = 9)
PARAMETER (m = mmax - 1, n = nmax - 1)
REAL x(0:mmax, 0:nmax)
...
DO k = 1, 300    ! iterations
  err1 = 0.0
  DO j = 1, n    ! columns
    DO i = 1, m  ! rows
      temp = 0.25 * (x(i,j-1) + x(i-1,j) + x(i+1,j) + x(i,j+1)) - x(i,j)
      x(i,j) = x(i,j) + omega * temp
      IF (abs(temp) > err1) err1 = abs(temp)
    ENDDO
  ENDDO
  IF (err1 <= eps) exit
ENDDO
...
END
```



# The Red-Black SOR Method



# The Sequential Algorithm

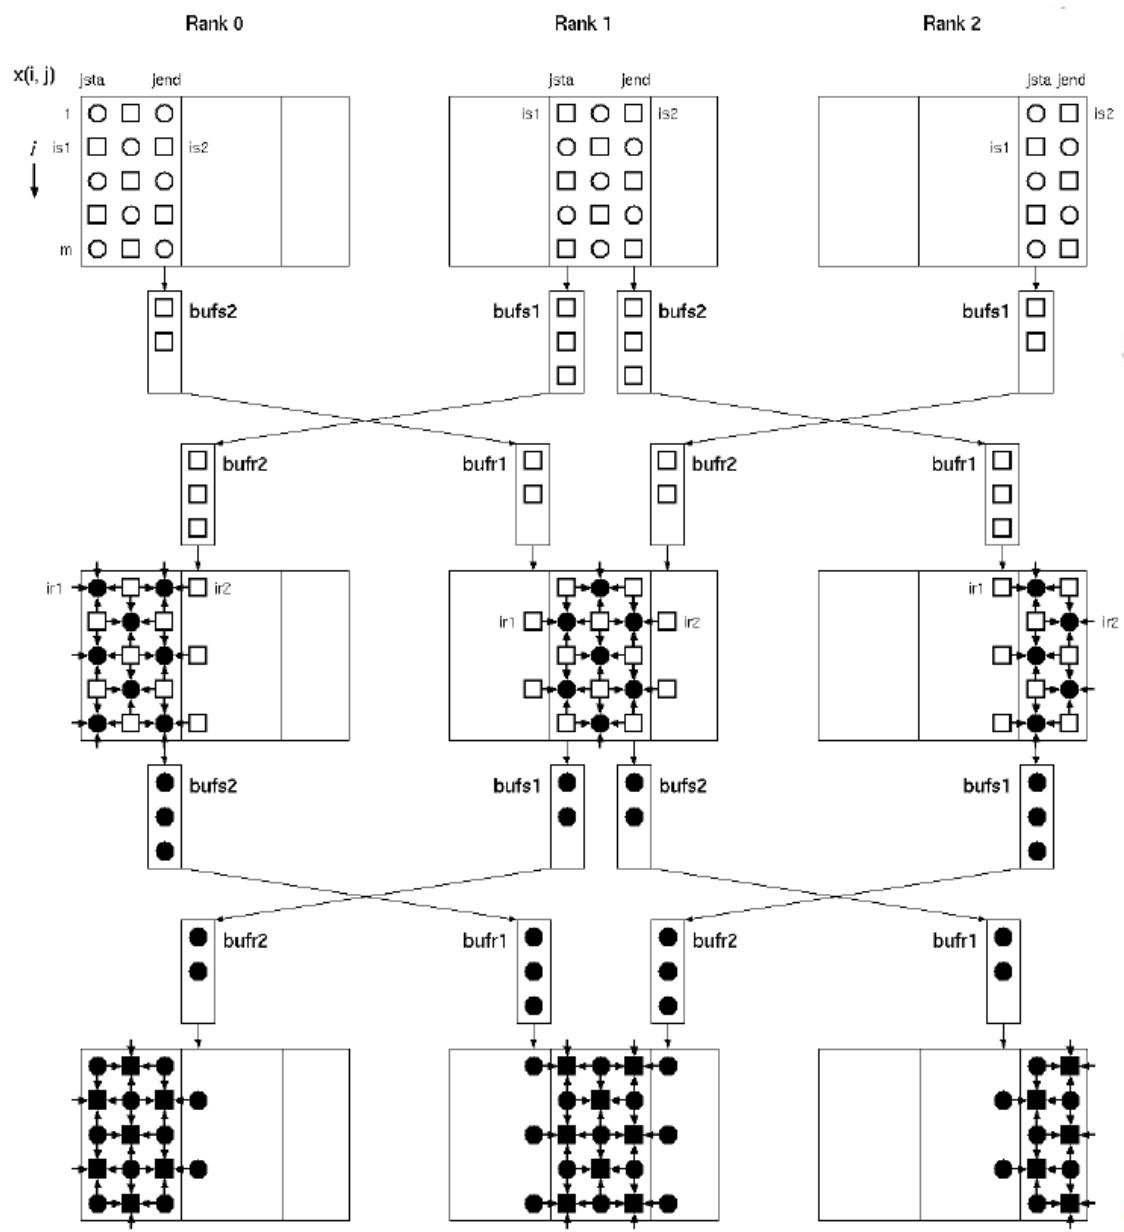
```

DO k = 1, 300    ! iterations
  err1 = 0.0
  DO j = 1, n    ! columns
    DO i = 1 + MOD(j+1,2), m, 2    ! circles on rows
      temp = 0.25 * (x(i,j-1) + x(i-1,j) + x(i+1,j) + x(i,j+1)) - x(i,j)
      x(i,j) = x(i,j) + omega * temp
      IF (abs(temp) > err1) err1 = abs(temp)
    ENDDO
  ENDDO
  DO j = 1, n    ! columns
    DO i = 1 + MOD(j,2), m, 2    ! squares on rows
      temp = 0.25 * (x(i,j-1) + x(i-1,j) + x(i+1,j) + x(i,j+1)) - x(i,j)
      x(i,j) = x(i,j) + omega * temp
      IF (abs(temp) > err1) err1 = abs(temp)
    ENDDO
  ENDDO
  IF (err1 <= eps) exit
ENDDO

```








● ■ : Updated in the current iteration  
○ □ : Not updated yet in the current iteration

# The Parallel Code

- Go to the ex4/version1 directory
- Look at main.f90 and grid.f90
- The data is distributed among the processes
- It is not clear whether the local block of data starts with a circle or a square
  - The variables **even** and **odd** keep track of whether the global column number of the first column in the local data is odd or even



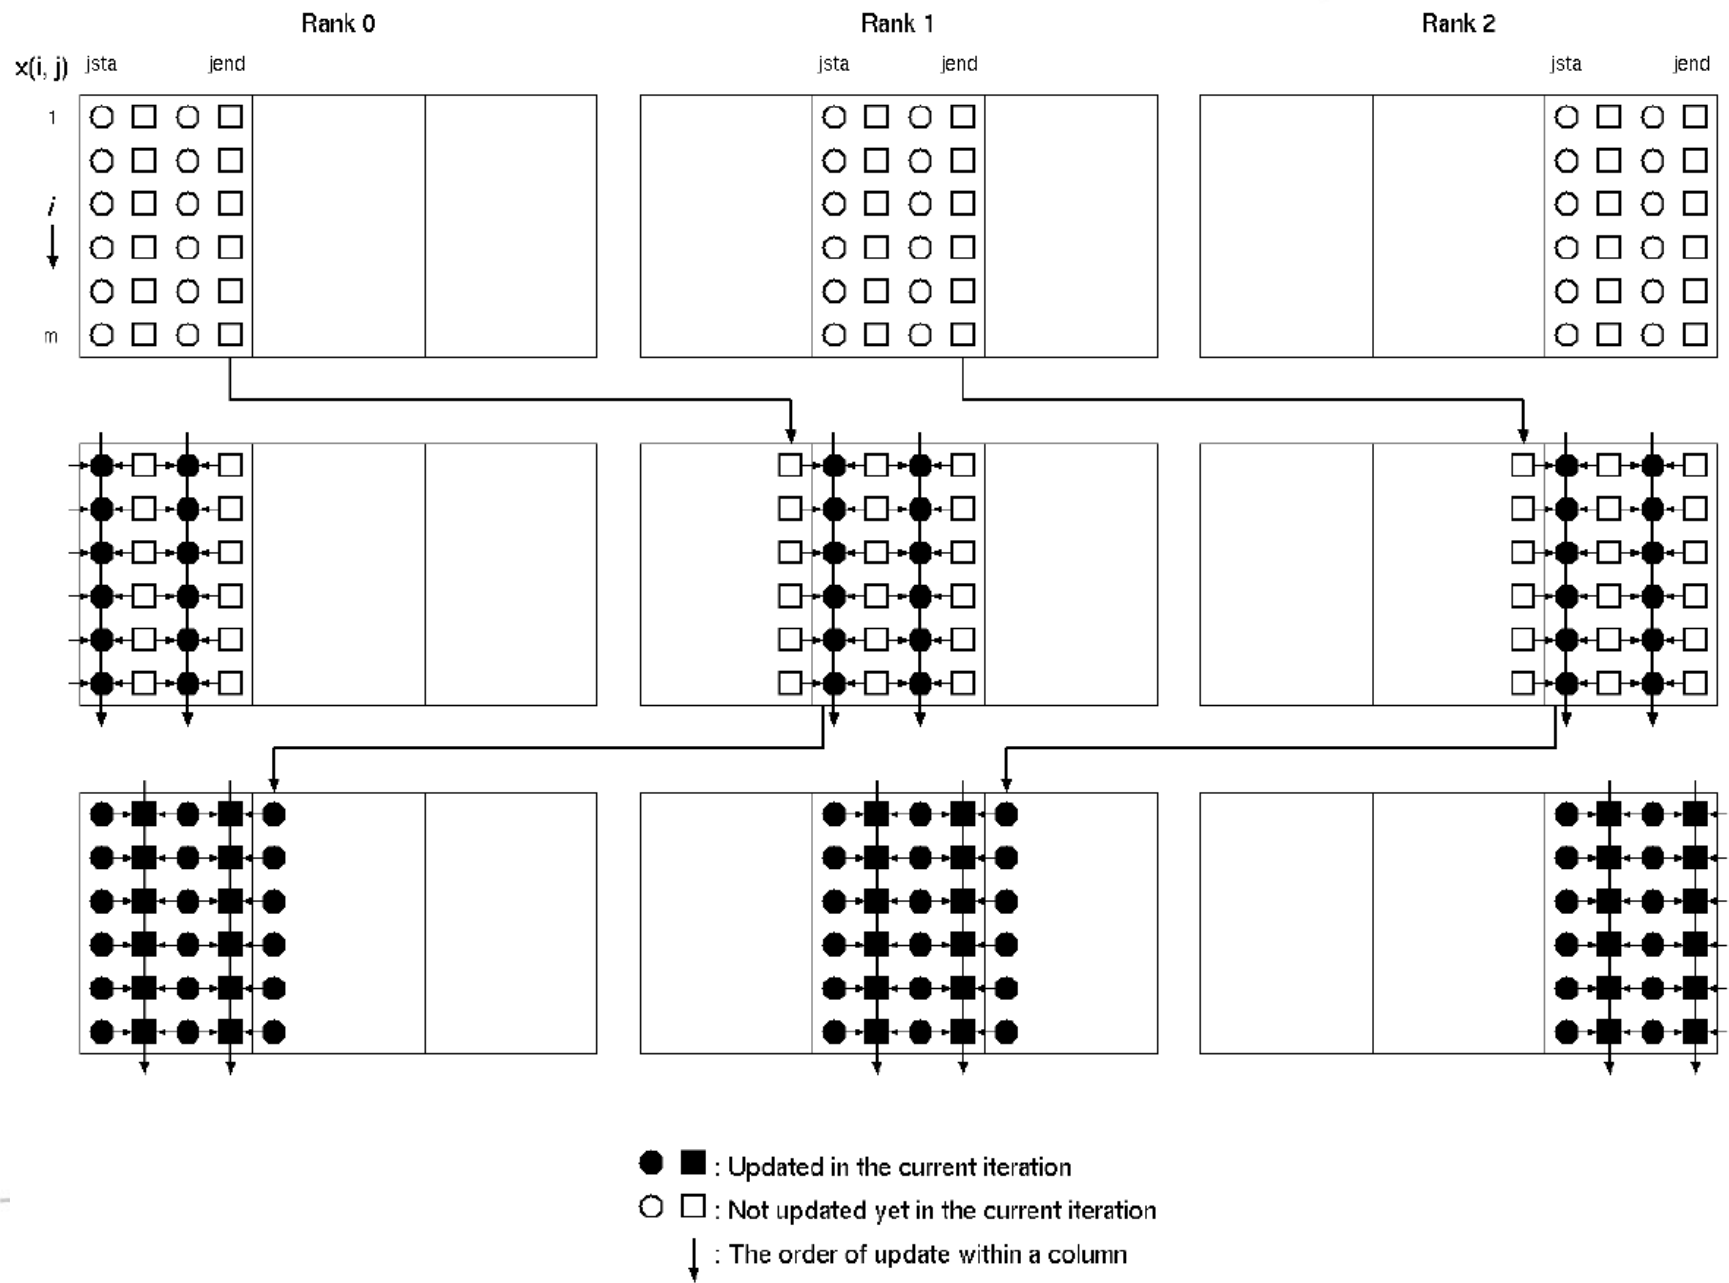
# The Zebra SOR Method



# The Sequential Algorithm

```
DO k = 1, 300
  err1 = 0.0
  DO j = 1, n, 2
    DO i = 1, m
      Update x(i,j) and err1
    ENDDO
  ENDDO
  DO j = 2, n, 2
    DO i = 1, m
      Update x(i,j) and err1
    ENDDO
  ENDDO
  IF (err1 <= eps) EXIT
ENDDO
```





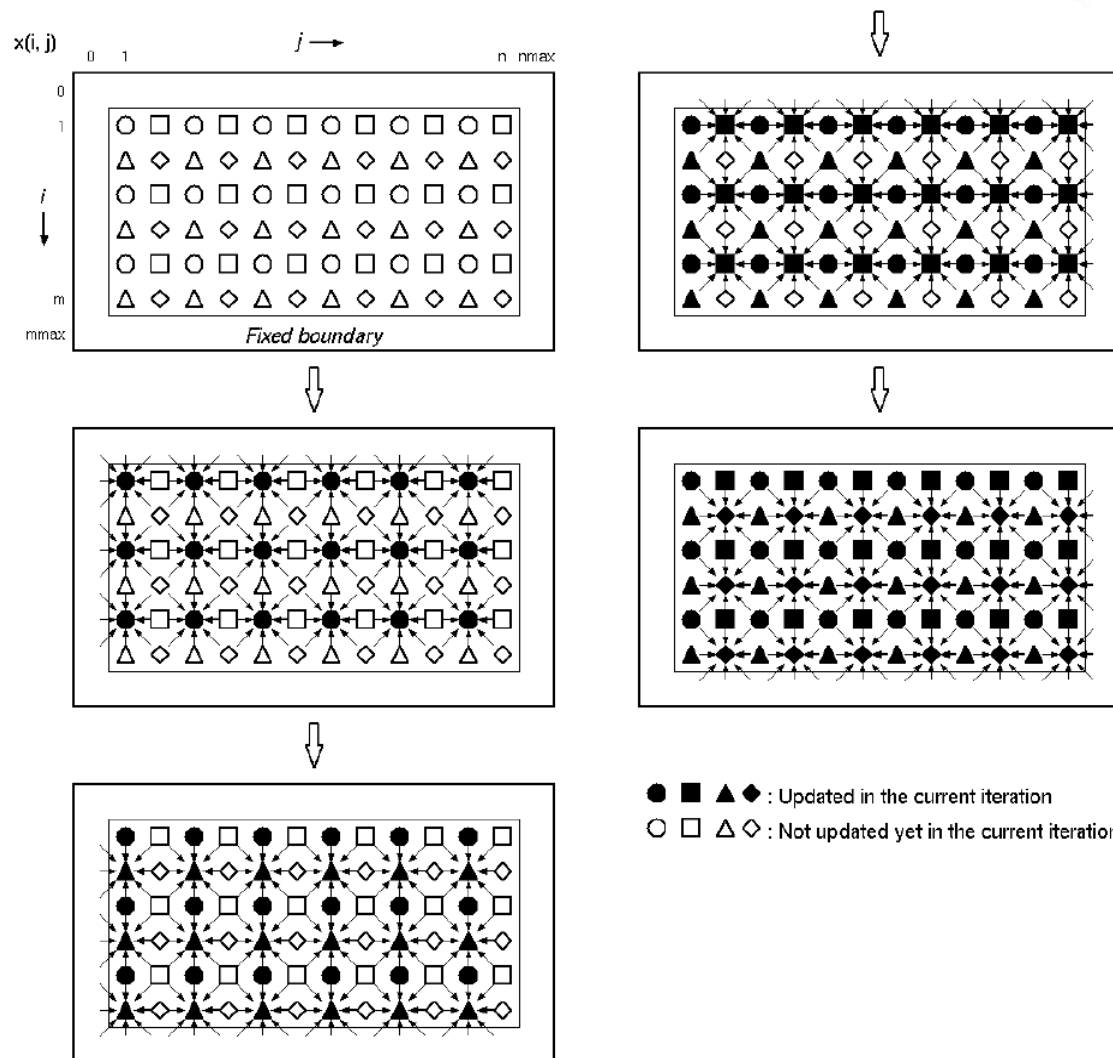
n  
of  
gy

# The Parallel Code

- Go to the ex4/version2 directory
- Look at main.f90 and grid.f90



# The Four-Colour SOR Method



# The Sequential Algorithm

```

PROGRAM fourcolor
PARAMETER (mmax = ..., nmax = ...)
PARAMETER (m = mmax - 1, n = nmax - 1)
DIMENSION x(0:mmax, 0:nmax)
...
DO k = 1, 300
  err1 = 0.0
  DO j = 1, n, 2
    DO i = 1, m, 2
      Update x(i,j) and err1
    ENDDO
  ENDDO
DO j = 1, n, 2
  DO i = 2, m, 2
    Update x(i,j) and err1
  ENDDO
ENDDO

```

```

DO j = 2, n, 2
  DO i = 1, m, 2
    Update x(i,j) and err1
  ENDDO
ENDDO
DO j = 2, n, 2
  DO i = 2, m, 2
    Update x(i,j) and err1
  ENDDO
ENDDO
IF (err1 <= eps) EXIT
ENDDO
...
END

```

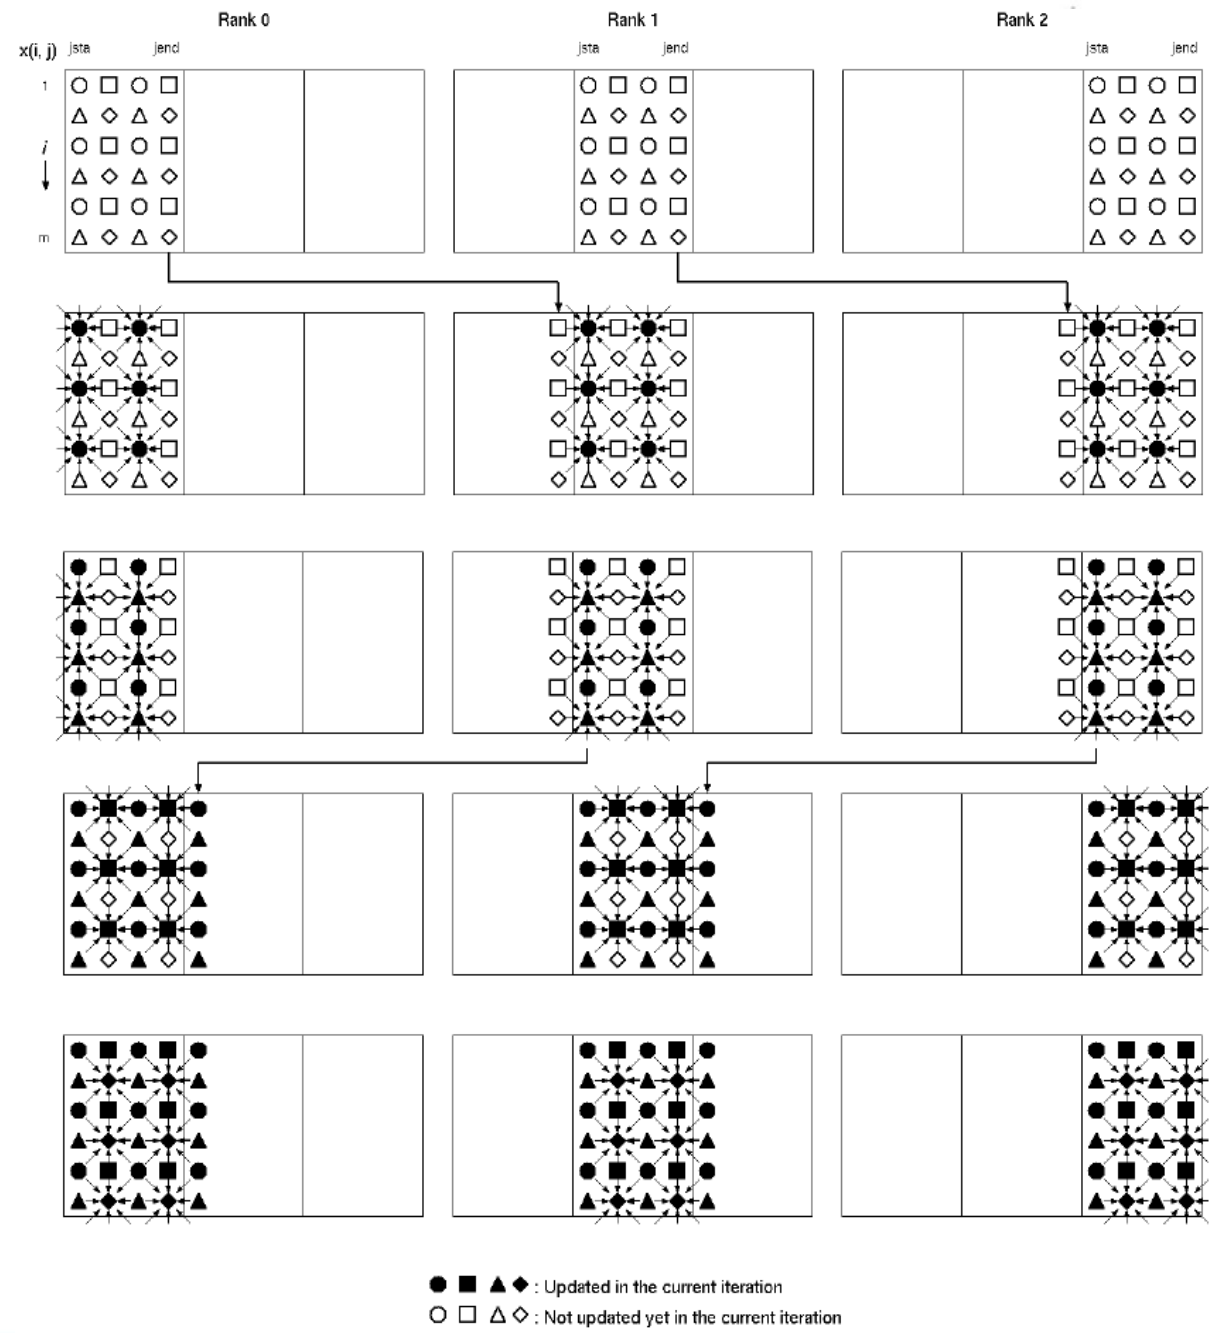




# 8 Neighbours are Involved

```
DO j = 1, n, 2
  DO i = 1, m, 2
    temp = 0.125 * (x(i, j-1) + x(i-1,j) + x(i+1,j) + x(i, j+1) + &
      x(i-1,j-1) + x(i+1,j-1) + x(i-1,j+1) + x(i+1,j+1)) - x(i,j)
    x(i,j) = x(i,j) + omega * temp
    IF (abs(temp) > err1) err1 = abs(temp)
  ENDDO
ENDDO
```





# The Parallel Code

- Go to the ex4/version3 directory
- Look at main.f90 and grid.f90




# The Monte Carlo Method

- Example 5 in the source code
- A random walk in 2D
- 100,000 particles
- 10 steps



# A Sample Trajectory



**NTNU – Trondheim**  
Norwegian University of  
Science and Technology

# The Sequential Algorithm

```
PROGRAM main
PARAMETER (n = 100000)
INTEGER itotal(0:9)
REAL seed
pi = 3.1415926
DO i = 0, 9
    itotal(i) = 0
ENDDO
seed = 0.5
CALL srand(seed)
```

```
DO i = 1, n
    x = 0.0
    y = 0.0
    DO istep = 1, 10
        angle = 2.0 * pi * rand()
        x = x + cos(angle)
        y = y + sin(angle)
    ENDDO
    itemp = sqrt(x**2 + y**2)
    itotal(itemp) = itotal(itemp) + 1
ENDDO
PRINT *, 'total =', itotal
END
```



# The Parallel Code

- Go to the directory ex5
- The sequential version is in directory version1
- The parallel version is in directory version2



# Molecular Dynamics

- Example 6 in the source code
- N particles interact in 1 dimension
- The force on particle i from particle j is given by

$$f_{ij} = 1/(x_j - x_i)$$


- The law of action and reaction applies:

$$f_{ij} = -f_{ji}$$





# Forces in 1D



# Forces Acting on 7 Particles

$f(1) =$		$+f_{12}$	$+f_{13}$	$+f_{14}$	$+f_{15}$	$+f_{16}$	$+f_{17}$
$f(2) =$	$-f_{12}$		$+f_{23}$	$+f_{24}$	$+f_{25}$	$+f_{26}$	$+f_{27}$
$f(3) =$	$-f_{13}$	$-f_{23}$		$+f_{34}$	$+f_{35}$	$+f_{36}$	$+f_{37}$
$f(4) =$	$-f_{14}$	$-f_{24}$	$-f_{34}$		$+f_{45}$	$+f_{46}$	$+f_{47}$
$f(5) =$	$-f_{15}$	$-f_{25}$	$-f_{35}$	$-f_{45}$		$+f_{56}$	$+f_{57}$
$f(6) =$	$-f_{16}$	$-f_{26}$	$-f_{36}$	$-f_{46}$	$-f_{56}$		$+f_{67}$
$f(7) =$	$-f_{17}$	$-f_{27}$	$-f_{37}$	$-f_{47}$	$-f_{57}$	$-f_{67}$	



# The Sequential Algorithm

```
PARAMETER (n = ...)
REAL f(n), x(n)
...
DO itime = 1, 100
  DO i = 1, n
    f(i) = 0.0
  ENDDO
  DO i = 1, n-1
    DO j = i+1, n
      fij = 1.0 / (x(j)-x(i))
      f(i) = f(i) + fij
      f(j) = f(j) - fij
    ENDDO
  ENDDO
  DO i = 1, n
    x(i) = x(i) + f(i)
  ENDDO
ENDDO
```



# Two Parallelisation Methods

- Cyclic distribution of the outer loop
- Cyclic distribution of the inner loop

```
DO i = 1, n-1
  DO j = i+1, n
    fij = 1.0 / (x(j)-x(i))
    f(i) = f(i) + fij
    f(j) = f(j) - fij
  ENDDO
ENDDO
```



Process 0

f(1) =		+f12	+f13	+f14	+f15	+f16	+f17
f(2) =	-f12						
f(3) =	-f13						
f(4) =	-f14			+f45	+f46	+f47	
f(5) =	-f15		-f45				
f(6) =	-f16		-f46				
f(7) =	-f17		-f47				

Process 1

f(1) =			+f23	+f24	+f25	+f26	+f27
f(2) =							
f(3) =		-f23					
f(4) =		-f24					
f(5) =		-f25		+f56	+f57		
f(6) =		-f26		-f56			
f(7) =		-f27		-f57			

Process 2

f(1) =							
f(2) =							
f(3) =			+f34	+f35	+f36	+f37	
f(4) =			-f34				
f(5) =			-f35				
f(6) =			-f36			+f67	
f(7) =			-f37			-f67	

MPI\_ALLREDUCE

All Processes

ff(1) =		+f12	+f13	+f14	+f15	+f16	+f17
ff(2) =	-f12		+f23	+f24	+f25	+f26	+f27
ff(3) =	-f13	-f23		+f34	+f35	+f36	+f37
ff(4) =	-f14	-f24	-f34		+f45	+f46	+f47
ff(5) =	-f15	-f25	-f35	-f45		+f56	+f57
ff(6) =	-f16	-f26	-f36	-f46	-f56		+f67
ff(7) =	-f17	-f27	-f37	-f47	-f57	-f67	

Process 0

f(1) =		+f12		+f15		
f(2) =	-f12		+f23		+f25	
f(3) =		-f23		+f35		
f(4) =				+f45		
f(5) =	-f15		-f35	-f45		+f56
f(6) =		-f26		-f56		
f(7) =						

Process 1

f(1) =		+f13		+f16		
f(2) =			+f24		+f27	
f(3) =	-f13			+f35		
f(4) =		-f24		+f45		
f(5) =					+f57	
f(6) =	-f16		-f36	-f46		
f(7) =		-f27		-f57		

Process 2

f(1) =		+f14		+f17		
f(2) =			+f25			
f(3) =			+f34		+f37	
f(4) =	-f14		-f34		+f47	
f(5) =		-f25				
f(6) =					+f57	
f(7) =	-f17		-f37	-f47	-f67	

MPI\_ALLREDUCE

All Processes

ff(1) =	+f12	+f13	+f14	+f15	+f16	+f17
ff(2) =	-f12	+f23	+f24	+f25	+f26	+f27
ff(3) =	-f13	-f23	+f34	+f35	+f36	+f37
ff(4) =	-f14	-f24	-f34	+f45	+f46	+f47
ff(5) =	-f15	-f25	-f35	-f45	+f56	+f57
ff(6) =	-f16	-f26	-f36	-f46	-f56	+f67
ff(7) =	-f17	-f27	-f37	-f47	-f57	-f67

# MPMD Models

- Example 7 in the source code
- Multiple Programs Multiple Data
- Different programs run in parallel and communicate with each other



# MPMD Model

## Process 0

```
PROGRAM fluid
INCLUDE 'mpif.h'
...
CALL MPI_INIT
CALL MPI_COMM_SIZE
CALL MPI_COMM_RANK
...
DO itime = 1, n
```

Computation of  
Fluid Dynamics


```
CALL MPI_SEND
CALL MPI_RECV
ENDDO
...
END
```

## Process 1

```
PROGRAM struct
INCLUDE 'mpif.h'
...
CALL MPI_INIT
CALL MPI_COMM_SIZE
CALL MPI_COMM_RANK
...
DO itime = 1, n
```

Computation of  
Structural Analysis

```
CALL MPI_RECV
CALL MPI_SEND
ENDDO
...
END
```





# Master/Worker Programs

- The master coordinates the execution of all the other processes
- The master has a list of jobs that must be processed
- Suitable if:
  - The processing time varies greatly from job to job
  - Neither block nor cyclic distribution gives a good load balancing
  - A heterogeneous environment where the performance of the machines is not uniform



# More Information

- All examples are based on:
  - <http://www.redbooks.ibm.com/abstracts/sg245380.html>
- Our Web-site:
  - <http://www.hpc.ntnu.no/>
- Send an e-mail to:
  - [support-ntnu@notur.no](mailto:support-ntnu@notur.no)
  - [support-kongull@hpc.ntnu.no](mailto:support-kongull@hpc.ntnu.no)

